


FIGURE 2

An agent run has three outcomes, not two

10 replicate runs per question (claude-sonnet-4-5, temperature 1.0); correctness by gpt-5 majority vote.



BixBench: 3 capsules covering 14 questions, 10 replicate runs each = 140 runs.