

Scored failures: striped = benchmark's defect, crosshatched = agent's

Fig 2's states, with every non-correct run attributed by re-deriving both the answer key and the agent's answer from the raw capsule data (env/REVIEW_REPORT.md). Of 80 answered-incorrect runs, 74 trace to the key or the question, 6 to the agent; of 10 no-answer runs, 9 follow the harness's own broken install instruction.



gpt-5 majority-vote grading. bix-8-q5's print-floor key is absent only because gpt-5 grades $p = 8.1 \times 10^{-194}$ as satisfying " $p < 2.2 \times 10^{-16}$ ".