

Most scored failures are fixable benchmark-side; hatching marks the agent's share

Fig 2's states, with every non-correct run attributed by re-deriving both the answer key and the agent's answer from the raw capsule data (env/REVIEW_REPORT.md). Unmarked red and gray segments trace to the benchmark (answer key, question wording, or the harness's own install instruction): 74 of 80 incorrect runs and 9 of 10 no-answer runs. The hatched segments — 6 incorrect, 1 no-answer — are the agent's.



gpt-5 majority-vote grading. bix-8-q5's print-floor key is absent only because gpt-5 grades $p = 8.1e-194$ as satisfying " $p < 2.2e-16$ ".