


Grader disagreement follows model generation, not model family

Fraction of answered runs graded correct per question (majority of 10 grading replicates per grader).
gpt-5 and claude-sonnet-4-5 — different families, current generation — agree on all 130 answers (nested dots);
gpt-4o — BixBench's default grader, one generation behind its own family-mate — flips whole questions.



Same 130 answers under all three graders.