


Given an answer, correctness is a property of the question, not luck

Answered runs only. Each dot is a question's estimated underlying success probability (posterior median, hierarchical model); the band is 95% uncertainty about that probability — 10 runs cannot pin it exactly even when all 10 agree.



Bernoulli GLMMs, question nested in capsule; gpt-5 majority-vote grading. Question-level SD (stage 2): 3.2 log-odds.