


Most incorrect answers trace to the benchmark; hatching marks the agent's share

Every incorrect run attributed by re-deriving both the answer key and the agent's answer from the raw capsule data (env/REVIEW_REPORT.md): 73 of 80 trace to the benchmark — the answer key or the question's wording. The 7 hatched runs, all in bix-26, are the agent's own.



gpt-5 majority-vote grading. bix-8-q5's print-floor key is absent only because gpt-5 grades $p = 8.1e-194$ as satisfying " $p < 2.2e-16$ ".