


FIGURE 4

Most incorrect answers trace to the benchmark

Every incorrect run attributed by re-deriving both the answer key and the agent's answer from the raw capsule data (env/REVIEW_REPORT.md).



gpt-5 majority-vote grading. bix-8-q5's print-floor key is absent only because gpt-5 grades $p = 8.1\text{e-}194$ as satisfying " $p < 2.2\text{e-}16$ ".